

# CP-2

February 3, 2024

## 0.1 HealthCare Capstone Project

### Importing Required libraries

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

### Importing Required Dataset

```
[3]: df = pd.read_csv('health care diabetes.csv')
```

### Getting shape of the dataset

```
[4]: df.shape
```

```
[4]: (768, 9)
```

### Getting info about datatype

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction              768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
```

```
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

### Checking top and bottom records

```
[6]: df.head(2)
```

```
[6]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6     148             72             35         0  33.6
1             1      85             66             29         0  26.6

   DiabetesPedigreeFunction  Age  Outcome
0                      0.627   50         1
1                      0.351   31         0
```

```
[7]: df.tail(2)
```

```
[7]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
766             1     126             60             0         0  30.1
767             1      93             70             31         0  30.4

   DiabetesPedigreeFunction  Age  Outcome
766                      0.349   47         1
767                      0.315   23         0
```

### Checking null values in dataset

```
[7]: df.isna().sum()
```

```
[7]: Pregnancies           0
Glucose                 0
BloodPressure           0
SkinThickness           0
Insulin                 0
BMI                     0
DiabetesPedigreeFunction 0
Age                     0
Outcome                 0
dtype: int64
```

### No null values in dataset

Values = 0 in feature columns (except pregnancies) means that value is missing and should be treated as null values


### 0.1.1 Checking '0' values in various columns with Visualization

```
[8]: (df.Glucose==0).sum()
```

```
[8]: 5
```

5 values are 0 in Glucose column

```
[9]: #Glucose
sns.histplot(data=df.Glucose,kde=True)
plt.title('Glucose level distribution of the patients')
plt.show()
plt.savefig('Glucose.png')
```



<Figure size 432x288 with 0 Axes>


Checking '0' values in BloodPressure columns

```
[10]: (df.BloodPressure==0).sum()
```

```
[10]: 35
```

35 values are 0 in BloodPressure column

```
[11]: #BloodPressure
sns.histplot(data=df.BloodPressure,kde=True)
plt.title('BloodPressure distribution of the patients')
plt.show()
plt.savefig('BloodPressure.png')
```



<Figure size 432x288 with 0 Axes>

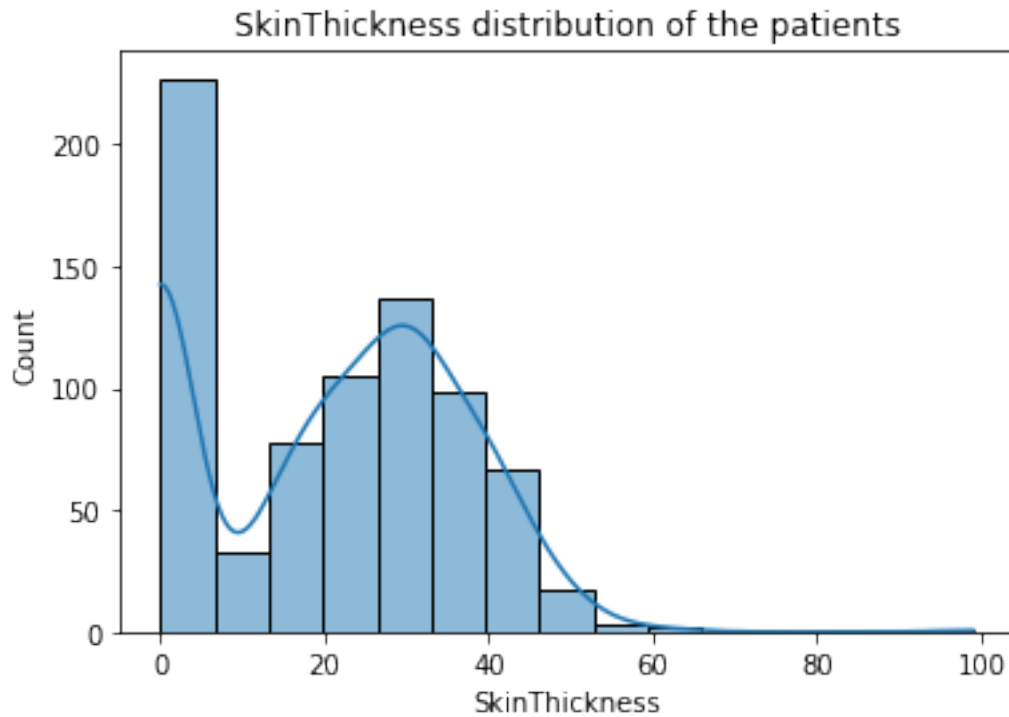
Checking '0' values in SkinThickness column

```
[12]: (df.SkinThickness==0).sum()
```

[12]: 227

227 values are 0 in SkinThickness column

```
[13]: #SkinThickness
sns.histplot(data=df.SkinThickness,kde=True)
plt.title('SkinThickness distribution of the patients')
plt.show()
plt.savefig('SkinThickness.png')
```



<Figure size 432x288 with 0 Axes>


Checking '0' values in Insulin column

```
[14]: (df.Insulin==0).sum()
```

```
[14]: 374
```

374 values are 0 in Insulin column

```
[15]: #Insulin
sns.histplot(data=df.Insulin,kde=True)
plt.title('Insulin distribution of the patients')
plt.show()
plt.savefig('Insulin.png')
```



<Figure size 432x288 with 0 Axes>

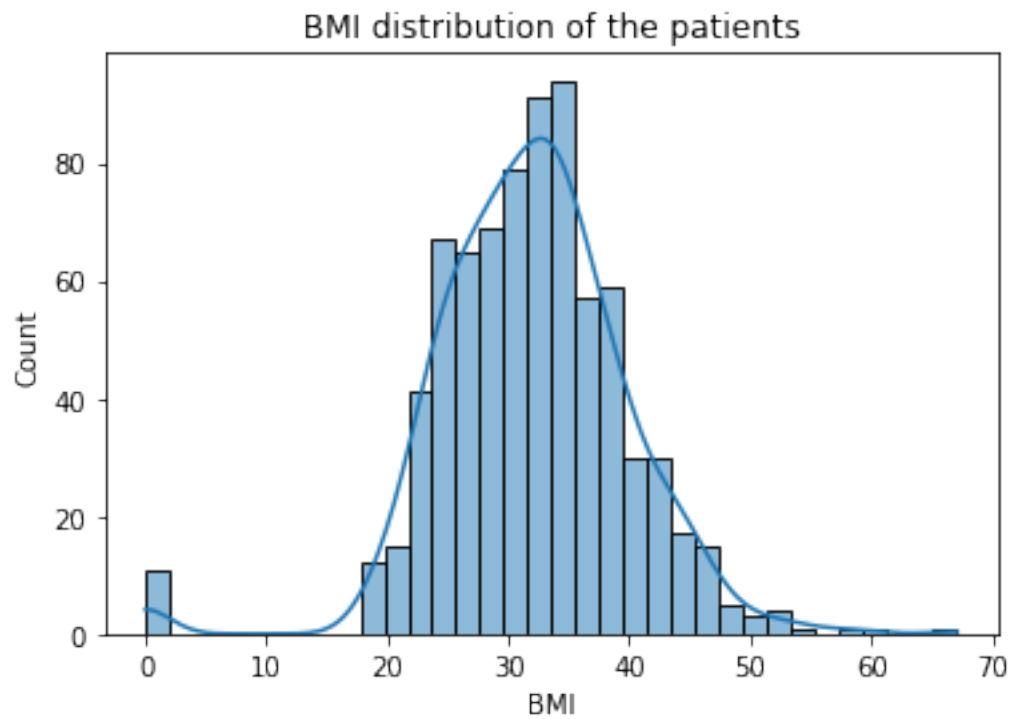
Checking '0' values in Insulin column

```
[16]: (df.BMI==0).sum()
```

```
[16]: 11
```

11 values are 0 in Insulin column

```
[17]: #BMI
sns.histplot(data=df.BMI,kde=True)
plt.title('BMI distribution of the patients')
plt.show()
plt.savefig('BMI.png')
```




<Figure size 432x288 with 0 Axes>

### 0.1.2 Checking outliers present in dataset

```
[18]: df.boxplot(figsize=(15,10))
```

```
[18]: <AxesSubplot: >
```



Insulin having huge outlier present as '0' which is missing values

### 0.1.3 Treating missing values

Importing simple imputer to replace 0 with the median of column

```
[19]: from sklearn.impute import SimpleImputer
impute=SimpleImputer(missing_values=0,strategy='median') #0 is treated as the
missing value
```

Creating new dataframe with imputed values

```
[20]: imputed=pd.DataFrame(impute.
fit_transform(df[['Glucose','BloodPressure','SkinThickness','Insulin','BMI']]),columns=('Glucose',
'BloodPressure', 'SkinThickness', 'Insulin', 'BMI'))
imputed
```

```
[20]:
```

	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	148.0	72.0	35.0	125.0	33.6
1	85.0	66.0	29.0	125.0	26.6
2	183.0	64.0	29.0	125.0	23.3
3	89.0	66.0	23.0	94.0	28.1
4	137.0	40.0	35.0	168.0	43.1
..	...	...	...	...	...
763	101.0	76.0	48.0	180.0	32.9
764	122.0	70.0	27.0	125.0	36.8




765	121.0	72.0	23.0	112.0	26.2
766	126.0	60.0	29.0	125.0	30.1
767	93.0	70.0	31.0	125.0	30.4

[768 rows x 5 columns]

comparing imputed and original data with histplot


```
[21]: plt.figure(figsize=(10,10))
plt.subplot(321)
sns.histplot(data=df.Glucose,kde=True)
plt.subplot(321)
sns.histplot(data=imputed.Glucose,color='r',kde=True)
```

[21]: <AxesSubplot: xlabel='Glucose', ylabel='Count'>




```
[22]: sns.histplot(data=df.BloodPressure,kde=True)
sns.histplot(data=imputed.BloodPressure,color='r',kde=True)
```

[22]: <AxesSubplot: xlabel='BloodPressure', ylabel='Count'>




```
[23]: sns.histplot(data=df.SkinThickness,kde=True)
      sns.histplot(data=imputed.SkinThickness,color='r',kde=True)
```

```
[23]: <AxesSubplot: xlabel='SkinThickness', ylabel='Count'>
```




```
[24]: sns.histplot(data=df.Insulin,kde=True)
      sns.histplot(data=imputed.Insulin,color='r',kde=True)
```

```
[24]: <AxesSubplot: xlabel='Insulin', ylabel='Count'>
```



```
[25]: sns.histplot(data=df.BMI,kde=True)
      sns.histplot(data=imputed.BMI,color='r',kde=True)
```

```
[25]: <AxesSubplot: xlabel='BMI', ylabel='Count'>
```



```
[26]: df.Glucose=imputed.Glucose
df.BloodPressure=imputed.BloodPressure
df.SkinThickness=imputed.SkinThickness
df.Insulin=imputed.Insulin
df.BMI=imputed.BMI
```

```
[27]: df
```

```
[27]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148.0	72.0	35.0	125.0	33.6	
1	1	85.0	66.0	29.0	125.0	26.6	
2	8	183.0	64.0	29.0	125.0	23.3	
3	1	89.0	66.0	23.0	94.0	28.1	
4	0	137.0	40.0	35.0	168.0	43.1	
..	...	...	...	...	...	...	
763	10	101.0	76.0	48.0	180.0	32.9	
764	2	122.0	70.0	27.0	125.0	36.8	
765	5	121.0	72.0	23.0	112.0	26.2	
766	1	126.0	60.0	29.0	125.0	30.1	
767	1	93.0	70.0	31.0	125.0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1

3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[28]: df[df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] == 0].
      count()
```

```
[28]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

Now, no null values are present

### Checking for any duplicate values in data

```
[29]: df[df.duplicated()]
```

```
[29]: Empty DataFrame
      Columns: [Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI,
      DiabetesPedigreeFunction, Age, Outcome]
      Index: []
```

```
[30]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   float64
2   BloodPressure          768 non-null   float64
3   SkinThickness          768 non-null   float64
4   Insulin                768 non-null   float64
```

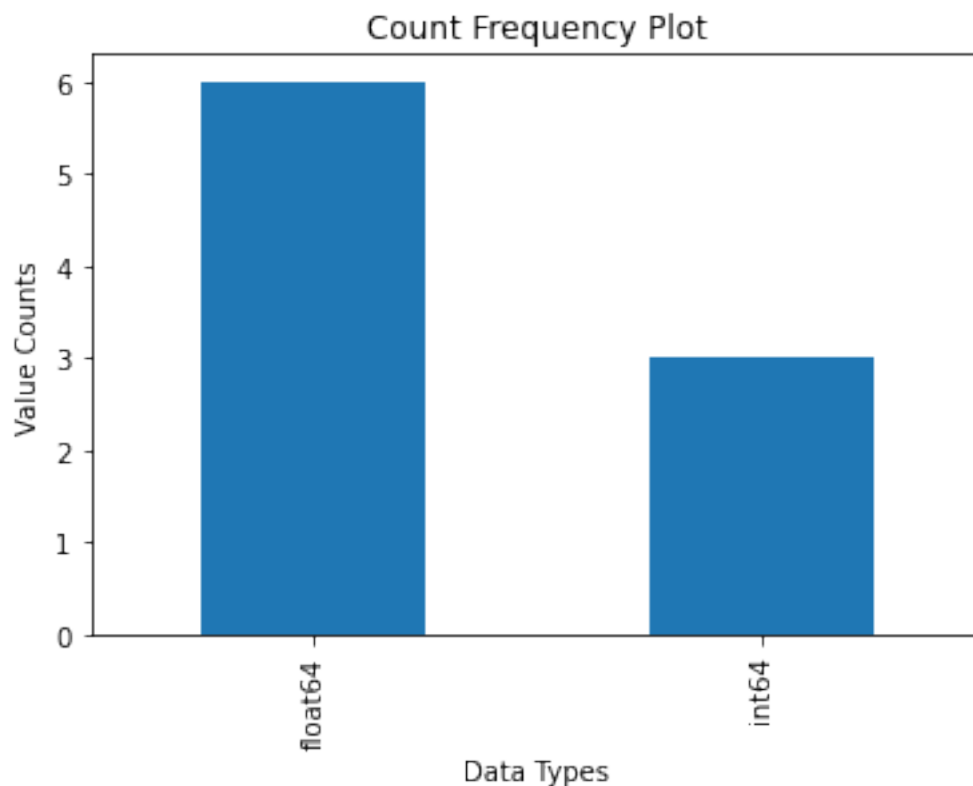
```
5   BMI                                768 non-null    float64
6   DiabetesPedigreeFunction          768 non-null    float64
7   Age                               768 non-null    int64
8   Outcome                           768 non-null    int64
dtypes: float64(6), int64(3)
memory usage: 54.1 KB
```

Plotting count frequency plot for describing the data types and the count of variables

```
[31]: df.dtypes.value_counts()
```

```
[31]: float64    6
      int64     3
      dtype: int64
```

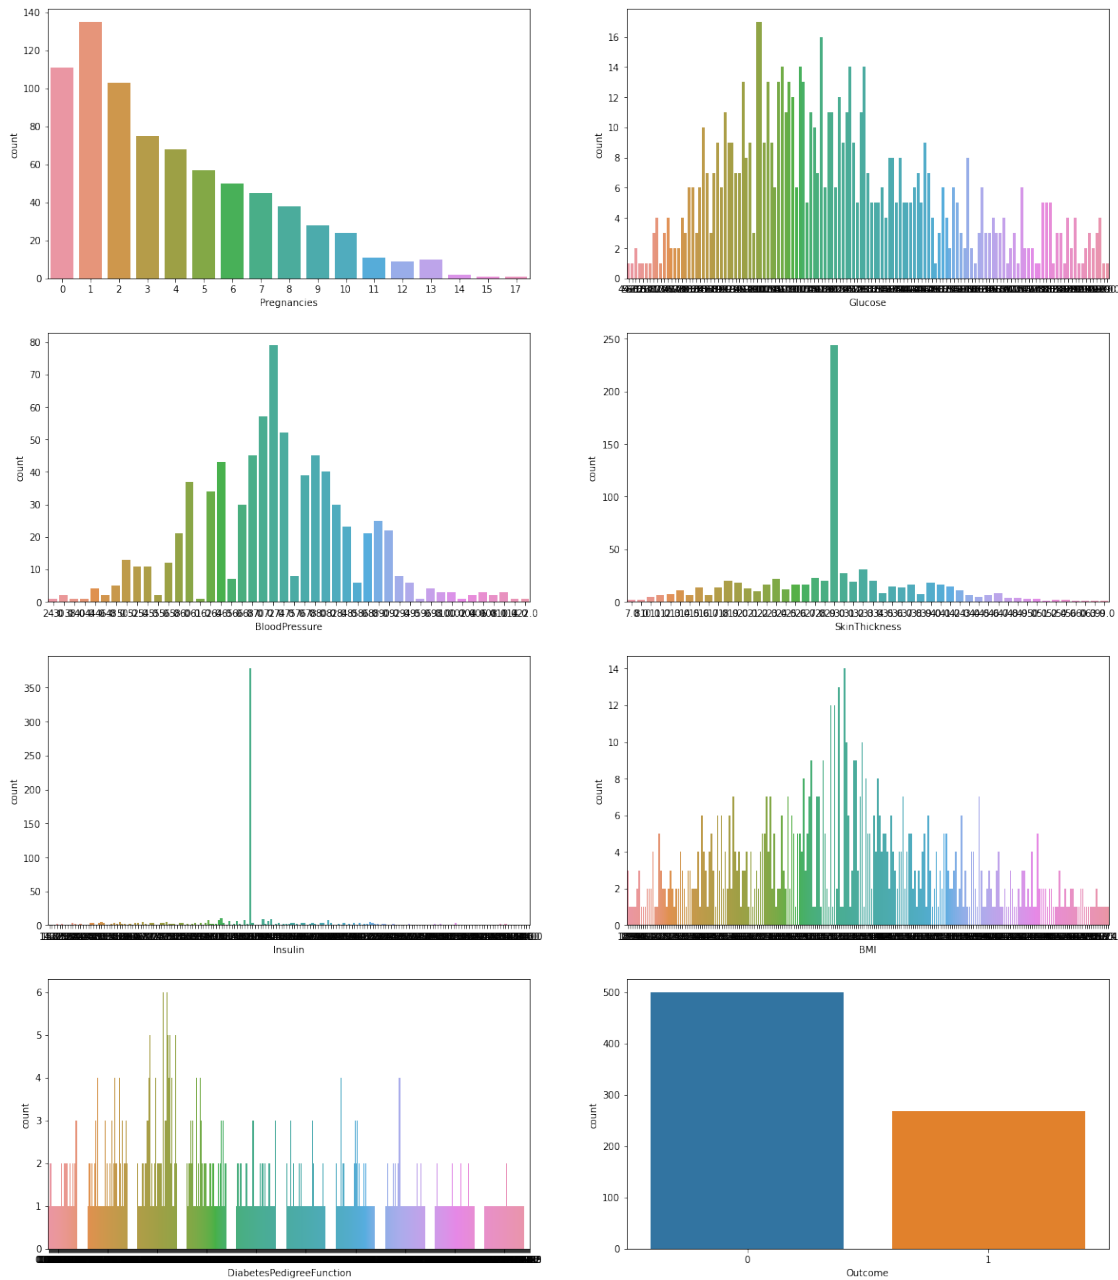
```
[32]: df1=df.dtypes.value_counts()
      df1.plot.bar()
      plt.title('Count Frequency Plot')
      plt.xlabel('Data Types')
      plt.ylabel('Value Counts')
      plt.show()
```



### Plotting countplot to show count of variables

```
[33]: plt.figure(figsize=(20,24))
plt.subplot(421)
sns.countplot(x=df.Pregnancies)
plt.subplot(422)
sns.countplot(x=df.Glucose)
plt.subplot(423)
sns.countplot(x=df.BloodPressure)
plt.subplot(424)
sns.countplot(x=df.SkinThickness)
plt.subplot(425)
sns.countplot(x=df.Insulin)
plt.subplot(426)
sns.countplot(x=df.BMI)
plt.subplot(427)
sns.countplot(x=df.DiabetesPedigreeFunction)
plt.subplot(428)
sns.countplot(x=df.Outcome)
```

```
[33]: <AxesSubplot: xlabel='Outcome', ylabel='count'>
```



Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action.

```
[34]: df['Outcome'].value_counts()
```

```
[34]: 0    500
      1    268
      Name: Outcome, dtype: int64
```



```
[35]: df['Outcome'].value_counts(normalize=True)
```

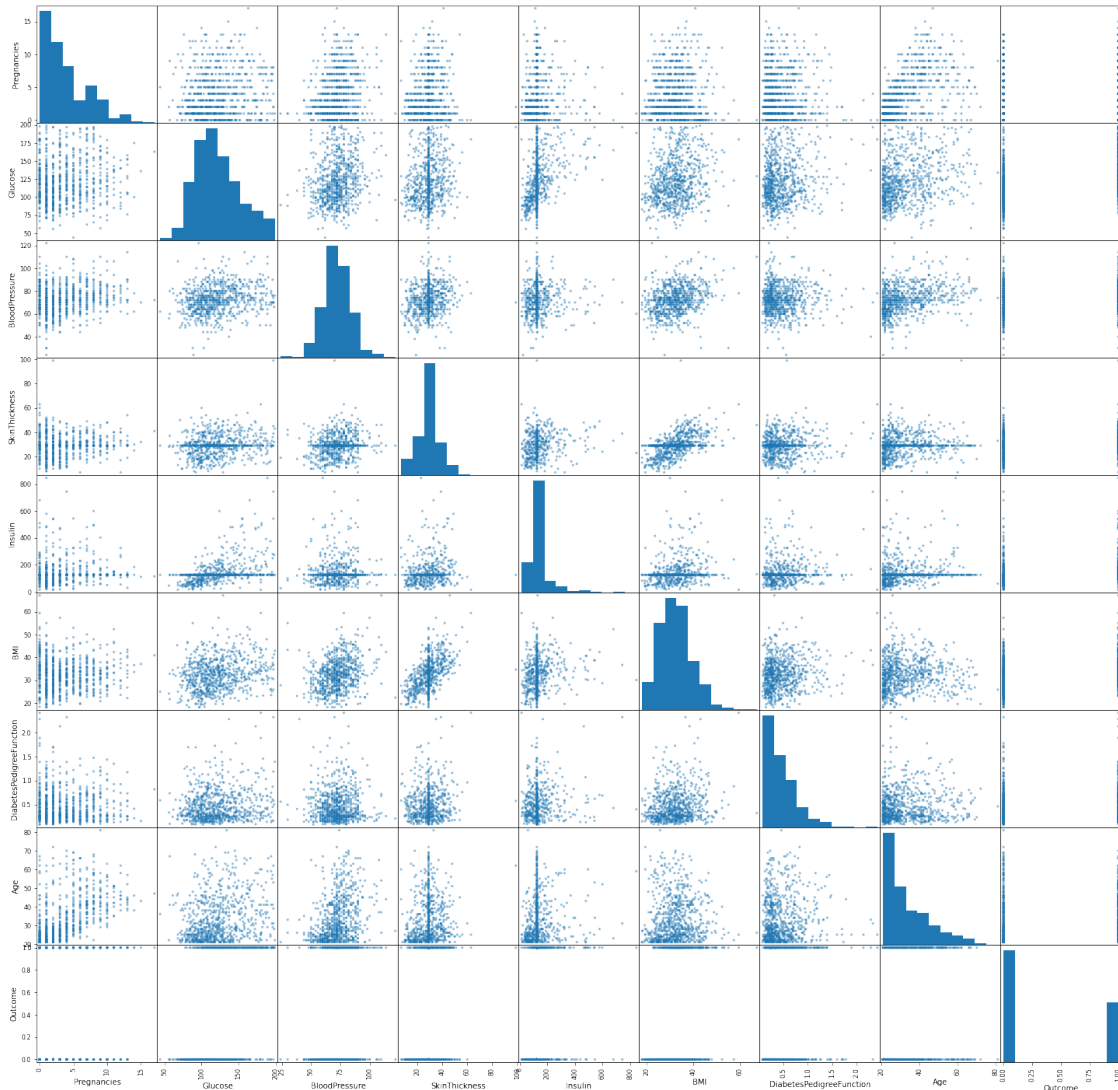
```
[35]: 0    0.651042  
      1    0.348958  
      Name: Outcome, dtype: float64
```

there are 500 negative or 0 and 268 positive or 1 instances hence we can say that the given dataset is imbalanced

In order to balance the data, we will add weights to each class.

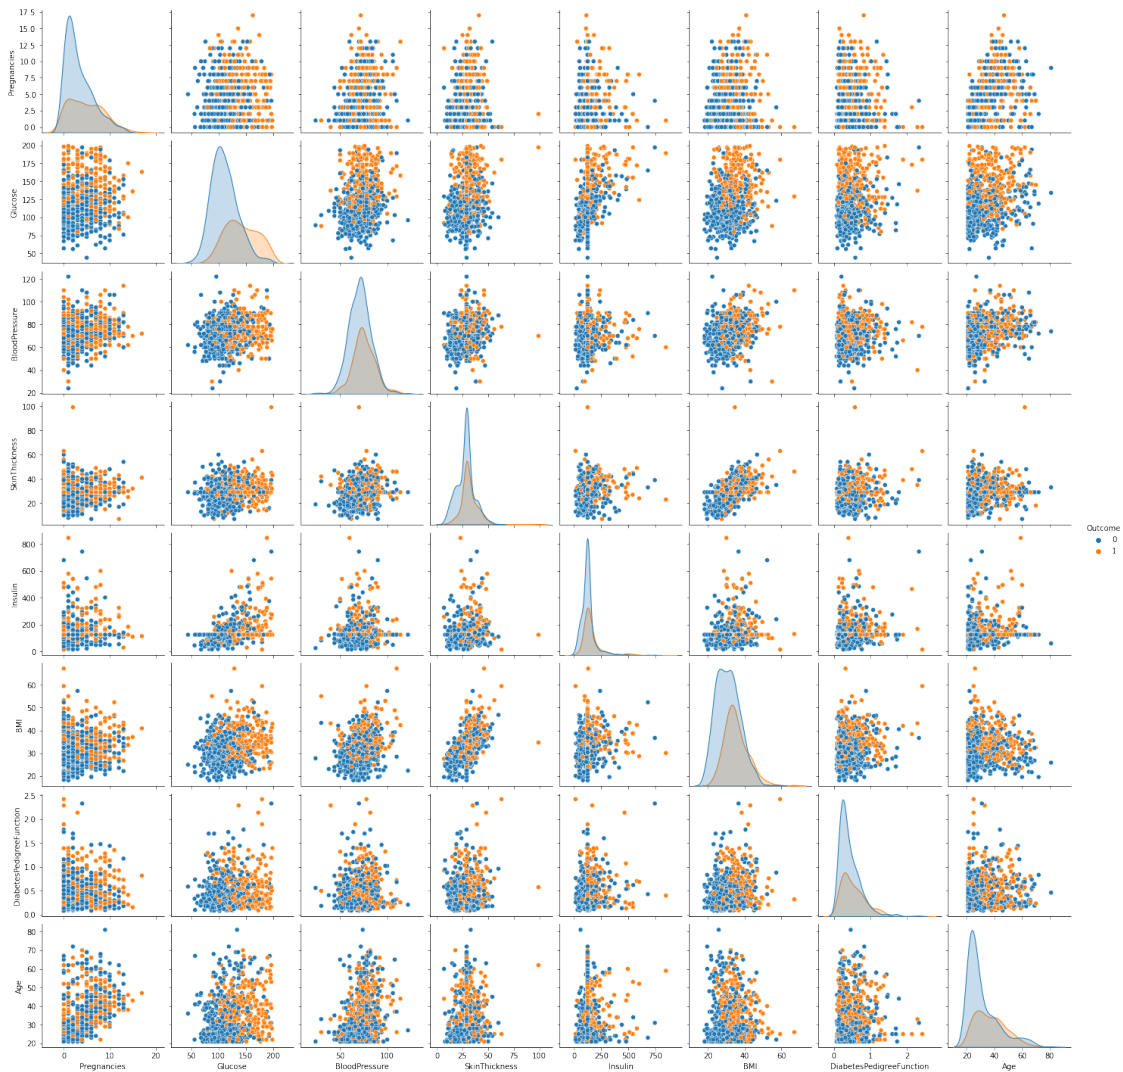
#### 0.1.4 Plotting scatter charts between the pair of variables to understand the relationships

```
[36]: from pandas.plotting import scatter_matrix  
      scatter = scatter_matrix(df, figsize=(25, 25))
```



```
[37]: sns.pairplot(df, hue = 'Outcome')
```

```
[37]: <seaborn.axisgrid.PairGrid at 0x7fb3e2256bf0>
```



### 0.1.5 Plotting heatmap for performing correlation analysis

Here, Value 1 represent the max correlation between variables Value 0 represent the no correlation between variables

```
[38]: df.corr()
```

```
[38]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.128213	0.208615	0.081770	
Glucose	0.128213	1.000000	0.218937	0.192615	
BloodPressure	0.208615	0.218937	1.000000	0.191892	
SkinThickness	0.081770	0.192615	0.191892	1.000000	
Insulin	0.025047	0.419451	0.045363	0.155610	
BMI	0.021559	0.231049	0.281257	0.543205	
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378	0.102188	
Age	0.544341	0.266909	0.324915	0.126107	
Outcome	0.221898	0.492782	0.165723	0.214873	

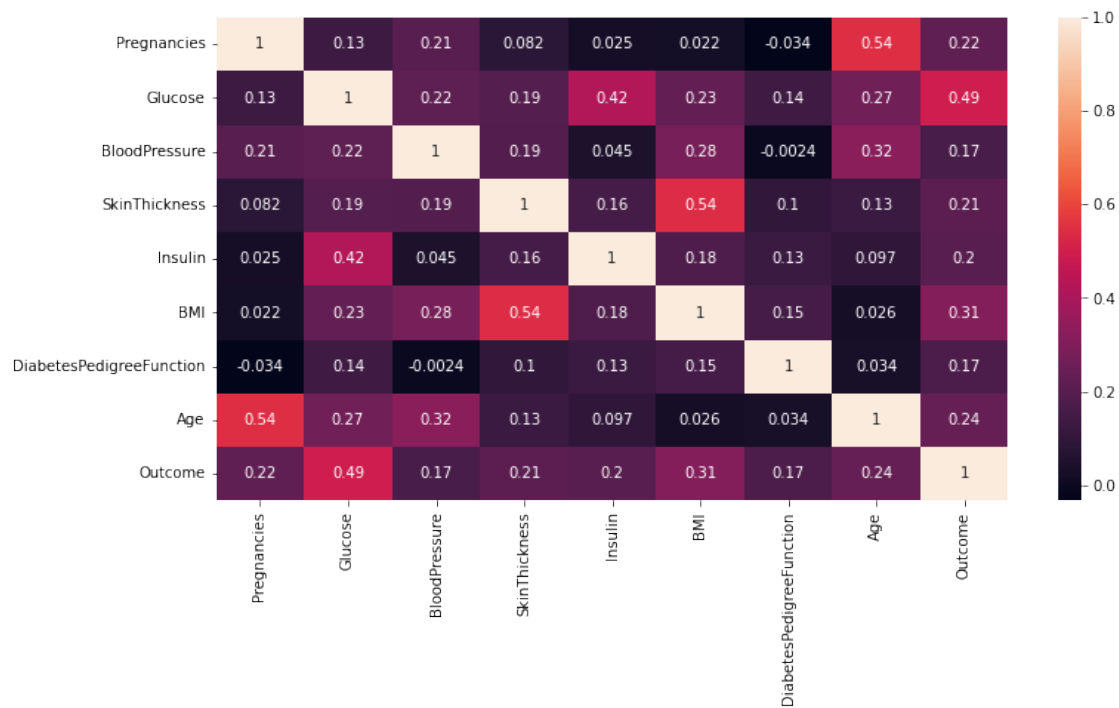
  

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	0.025047	0.021559	-0.033523	
Glucose	0.419451	0.231049	0.137327	
BloodPressure	0.045363	0.281257	-0.002378	
SkinThickness	0.155610	0.543205	0.102188	
Insulin	1.000000	0.180241	0.126503	
BMI	0.180241	1.000000	0.153438	
DiabetesPedigreeFunction	0.126503	0.153438	1.000000	
Age	0.097101	0.025597	0.033561	
Outcome	0.203790	0.312038	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.126107	0.214873
Insulin	0.097101	0.203790
BMI	0.025597	0.312038
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[39]: plt.figure(figsize = (12, 6))
heatmap = sns.heatmap(df.corr(), annot=True)
```



## 0.2 Data Modeling

Splitting the data into feature and Target variable

```
[71]: x = df.iloc[:, :-1]
      y = df.iloc[:, -1:]
      x,y
```

```
[71]: (
      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0                6    148.0           72.0          35.0    125.0  33.6
1                1     85.0           66.0          29.0    125.0  26.6
2                8    183.0           64.0          29.0    125.0  23.3
3                1     89.0           66.0          23.0     94.0  28.1
4                0    137.0           40.0          35.0    168.0  43.1
..            ...      ...           ...           ...      ...
763             10    101.0           76.0          48.0    180.0  32.9
764              2    122.0           70.0          27.0    125.0  36.8
765              5    121.0           72.0          23.0    112.0  26.2
766              1    126.0           60.0          29.0    125.0  30.1
767              1     93.0           70.0          31.0    125.0  30.4

      DiabetesPedigreeFunction  Age
0                        0.627   50
1                        0.351   31
2                        0.672   32
```

```

3          0.167    21
4          2.288    33
..          ...    ...
763        0.171    63
764        0.340    27
765        0.245    30
766        0.349    47
767        0.315    23

```

```
[768 rows x 8 columns],
```

```

    Outcome
0          1
1          0
2          1
3          0
4          1
..          ...
763        0
764        0
765        0
766        1
767        0

```

```
[768 rows x 1 columns])
```

### Performing Train Test split

```
[72]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=10,test_size=0.
↪2)
x_train,x_test,y_train,y_test
```

```
[72]: (   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
120          0    162.0          76.0          56.0    100.0  53.2
172          2     87.0          72.0          23.0    125.0  28.9
307          0    137.0          68.0          14.0    148.0  24.8
7           10    115.0          72.0          29.0    125.0  35.3
448          0    104.0          64.0          37.0     64.0  33.6
..          ...    ...          ...          ...    ...    ...
369          1    133.0         102.0          28.0    140.0  32.8
320          4    129.0          60.0          12.0    231.0  27.5
527          3    116.0          74.0          15.0    105.0  26.3
125          1     88.0          30.0          42.0     99.0  55.0
265          5     96.0          74.0          18.0     67.0  33.6

    DiabetesPedigreeFunction  Age
120                0.759    25
```

172	0.773	25
307	0.143	21
7	0.134	29
448	0.510	22
..	...	...
369	0.234	45
320	0.527	31
527	0.107	24
125	0.496	26
265	0.997	43

[614 rows x 8 columns],

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
568	4	154.0	72.0	29.0	126.0	31.3	
620	2	112.0	86.0	42.0	160.0	38.4	
456	1	135.0	54.0	29.0	125.0	26.7	
197	3	107.0	62.0	13.0	48.0	22.9	
714	3	102.0	74.0	29.0	125.0	29.5	
..	...	...	...	...	...	...	
264	4	123.0	62.0	29.0	125.0	32.0	
706	10	115.0	72.0	29.0	125.0	32.3	
194	8	85.0	55.0	20.0	125.0	24.4	
179	5	130.0	82.0	29.0	125.0	39.1	
514	3	99.0	54.0	19.0	86.0	25.6	

	DiabetesPedigreeFunction	Age
568	0.338	37
620	0.246	28
456	0.687	62
197	0.678	23
714	0.121	32
..	...	...
264	0.226	35
706	0.261	30
194	0.136	42
179	0.956	37
514	0.154	24

[154 rows x 8 columns],

	Outcome
120	1
172	0
307	0
7	0
448	1
..	...
369	1

```

320      0
527      0
125      1
265      0

[614 rows x 1 columns],
      Outcome
568      0
620      0
456      0
197      1
714      0
..      ...
264      1
706      1
194      0
179      1
514      0

[154 rows x 1 columns])

```

### 0.2.1 Logistic Regression Model

```

[73]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import confusion_matrix, classification_report
      lr_model = LogisticRegression()
      lr_model.fit(x_train, y_train)

```

```

[73]: LogisticRegression()

```

```

[74]: print('Accuracy on Train Data = ', lr_model.score(x_train, y_train))
      print('Accuracy on Test Data = ', lr_model.score(x_test, y_test))
      print('Confusion matrix :\n ', confusion_matrix(y, lr_model.predict(x)))
      print('Classification report :\n ', classification_report(y, lr_model.predict(x)))

```

```

Accuracy on Train Data = 0.7833876221498371
Accuracy on Test Data = 0.7272727272727273
Confusion matrix :
[[445  55]
 [120 148]]
Classification report :

```

	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.55	0.63	268
accuracy			0.77	768

macro avg	0.76	0.72	0.73	768
weighted avg	0.77	0.77	0.76	768

## 0.2.2 Logistic Regression Model ( With Added weights => 'Balanced' )

```
[75]: lrw_model=LogisticRegression(class_weight='balanced')
lrw_model.fit(x_train,y_train)
print('Accuracy on Train Data = ',lrw_model.score(x_train,y_train))
print('Accuracy on Test Data = ',lrw_model.score(x_test,y_test))
print('Confusion matrix : \n',confusion_matrix(y,lrw_model.predict(x)))
print('Classification report : \n',classification_report(y,lrw_model.
↪predict(x)))
```

Accuracy on Train Data = 0.747557003257329

Accuracy on Test Data = 0.7207792207792207

Confusion matrix :

[[377 123]

[ 75 193]]

Classification report :

	precision	recall	f1-score	support
0	0.83	0.75	0.79	500
1	0.61	0.72	0.66	268
accuracy			0.74	768
macro avg	0.72	0.74	0.73	768
weighted avg	0.76	0.74	0.75	768

## 0.2.3 Logistic Regression Model ( With Manual weights )

```
[76]: # for class 0, weighth = 1
# for class 1, weight = occurances of class 0 / occurances of class 1 = 500/267
↪= 1.872
lrwm_model=LogisticRegression(class_weight={0:1,1:1.872})
lrwm_model.fit(x_train,y_train)
print('Accuracy on Train Data = ',lrwm_model.score(x_train,y_train))
print('Accuracy on Test Data = ',lrwm_model.score(x_test,y_test))
print('Confusion matrix : \n',confusion_matrix(y,lrwm_model.predict(x)))
print('Classification report : \n',classification_report(y,lrwm_model.
↪predict(x)))
```

Accuracy on Train Data = 0.747557003257329

Accuracy on Test Data = 0.7467532467532467

Confusion matrix :

[[381 119]

[ 75 193]]



```

Classification report :
              precision    recall  f1-score   support

     0           0.84       0.76       0.80         500
     1           0.62       0.72       0.67         268

 accuracy                   0.75         768
 macro avg           0.73       0.74       0.73         768
 weighted avg        0.76       0.75       0.75         768

```

### 0.2.4 Decision Tree Model

```

[77]: from sklearn.tree import DecisionTreeClassifier
      dt_model=DecisionTreeClassifier()
      dt_model.fit(x_train,y_train)

```

```

[77]: DecisionTreeClassifier()

```

```

[78]: print('Accuracy on Train Data = ',dt_model.score(x_train,y_train))
      print('Accuracy on Test Data = ',dt_model.score(x_test,y_test))
      print('Confusion matrix : \n',confusion_matrix(y,dt_model.predict(x)))
      print('Classification report : \n',classification_report(y,dt_model.predict(x)))

```

```

Accuracy on Train Data =  1.0
Accuracy on Test Data =  0.7402597402597403
Confusion matrix :
[[483  17]
 [ 23 245]]
Classification report :
              precision    recall  f1-score   support

     0           0.95       0.97       0.96         500
     1           0.94       0.91       0.92         268

 accuracy                   0.95         768
 macro avg           0.94       0.94       0.94         768
 weighted avg        0.95       0.95       0.95         768

```

### 0.2.5 Random Forest Model

```

[79]: from sklearn.ensemble import RandomForestClassifier
      rf_model=RandomForestClassifier()
      rf_model.fit(x_train,y_train)

```

```

[79]: RandomForestClassifier()

```

```
[80]: print('Accuracy on Train Data = ',rf_model.score(x_train,y_train))
print('Accuracy on Test Data = ',rf_model.score(x_test,y_test))
print('Confusion matrix : \n',confusion_matrix(y,rf_model.predict(x)))
print('Classification report : \n',classification_report(y,rf_model.predict(x)))
```

```
Accuracy on Train Data = 1.0
Accuracy on Test Data = 0.7077922077922078
Confusion matrix :
[[484 16]
 [ 29 239]]
Classification report :
```

	precision	recall	f1-score	support
0	0.94	0.97	0.96	500
1	0.94	0.89	0.91	268
accuracy			0.94	768
macro avg	0.94	0.93	0.93	768
weighted avg	0.94	0.94	0.94	768

## 0.2.6 SVC Model

```
[83]: from sklearn.svm import SVC
svc_model=SVC(probability=True)
svc_model.fit(x_train,y_train)
```

```
[83]: SVC(probability=True)
```

```
[84]: print('Accuracy on Train Data = ',svc_model.score(x_train,y_train))
print('Accuracy on Test Data = ',svc_model.score(x_test,y_test))
print('Confusion matrix : \n',confusion_matrix(y,svc_model.predict(x)))
print('Classification report : \n',classification_report(y,svc_model.
↪predict(x)))
```

```
Accuracy on Train Data = 0.7785016286644951
Accuracy on Test Data = 0.6948051948051948
Confusion matrix :
[[457 43]
 [140 128]]
Classification report :
```

	precision	recall	f1-score	support
0	0.77	0.91	0.83	500
1	0.75	0.48	0.58	268
accuracy			0.76	768
macro avg	0.76	0.70	0.71	768

weighted avg	0.76	0.76	0.75	768
--------------	------	------	------	-----

## 0.2.7 KNN Model

```
[85]: from sklearn.neighbors import KNeighborsClassifier
      knn_model=KNeighborsClassifier()
      knn_model.fit(x_train,y_train)
```

```
[85]: KNeighborsClassifier()
```

```
[86]: print('Accuracy on Train Data = ',knn_model.score(x_train,y_train))
      print('Accuracy on Test Data = ',knn_model.score(x_test,y_test))
      print('Confusion matrix : \n',confusion_matrix(y,knn_model.predict(x)))
      print('Classification report : \n',classification_report(y,knn_model.
      ↪predict(x)))
```

Accuracy on Train Data = 0.8127035830618893

Accuracy on Test Data = 0.6883116883116883

Confusion matrix :

```
[[433  67]
```

```
 [ 96 172]]
```

Classification report :

	precision	recall	f1-score	support
0	0.82	0.87	0.84	500
1	0.72	0.64	0.68	268
accuracy			0.79	768
macro avg	0.77	0.75	0.76	768
weighted avg	0.78	0.79	0.78	768

## 0.2.8 AUC and ROC Curve fo various Model

```
[87]: from sklearn.metrics import roc_auc_score,roc_curve
      # getting prediction probabilities for different models
      lr_probs=lr_model.predict_proba(x)
      lrw_probs=lrw_model.predict_proba(x)
      lrwm_probs=lrwm_model.predict_proba(x)
      dt_probs=dt_model.predict_proba(x)
      rf_probs=rf_model.predict_proba(x)
      svc_probs=svc_model.predict_proba(x)
      knn_probs=knn_model.predict_proba(x)
      # We need pred probs only for outcome = 1
      lr_probs=lr_probs[:,1]
      lrw_probs=lrw_probs[:,1]
      lrwm_probs=lrwm_probs[:,1]
```

```
dt_probs=dt_probs[:,1]
rf_probs=rf_probs[:,1]
svc_probs=svc_probs[:,1]
knn_probs=knn_probs[:,1]
```

```
[88]: # getting AUC score
auc_lr=roc_auc_score(y,lr_probs)
auc_lrw=roc_auc_score(y,lrw_probs)
auc_lrwm=roc_auc_score(y,lrwm_probs)
auc_dt=roc_auc_score(y,dt_probs)
auc_rf=roc_auc_score(y,rf_probs)
auc_svc=roc_auc_score(y,svc_probs)
auc_knn=roc_auc_score(y,knn_probs)
print('\n AUC for different models : \n')
print('Logistic regression model without weight AUC : ',auc_lr)
print('Logistic Regression model with Balanced weight AUC : ',auc_lrw)
print('Logistic Regression model with manual weight AUC : ',auc_lrwm)
print('Decision Tree model AUC : ',auc_dt)
print('Random Forest model AUC : ',auc_rf)
print('SVC model AUC : ',auc_svc)
print('KNN model AUC : ',auc_knn)
```

AUC for different models :

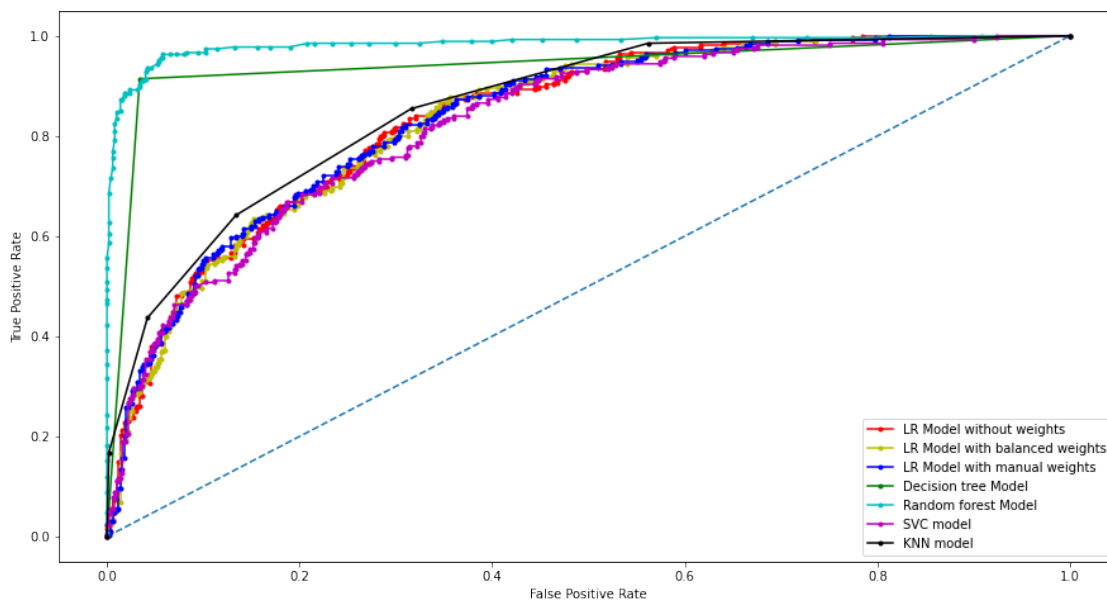
```
Logistic regression model without weight AUC :  0.8374477611940299
Logistic Regression model with Balanced weight AUC :  0.8353358208955224
Logistic Regression model with manual weight AUC :  0.8385820895522388
Decision Tree model AUC :  0.940089552238806
Random Forest model AUC :  0.9838768656716419
SVC model AUC :  0.8259701492537312
KNN model AUC :  0.8590186567164179
```

```
[89]: # calculating values for ROC curve
lr_fpr, lr_tpr, lr_thresholds = roc_curve(y, lr_probs)
lrw_fpr, lrw_tpr, lrw_thresholds = roc_curve(y, lrw_probs)
lrwm_fpr, lrwm_tpr, lrwm_thresholds = roc_curve(y, lrwm_probs)
dt_fpr, dt_tpr, dt_thresholds = roc_curve(y, dt_probs)
rf_fpr, rf_tpr, rf_thresholds = roc_curve(y, rf_probs)
svc_fpr, svc_tpr, svc_thresholds = roc_curve(y, svc_probs)
knn_fpr, knn_tpr, knn_thresholds = roc_curve(y, knn_probs)
```

```
[90]: # plotting for random predictor model
print('\nROC curves for various models \n')
plt.figure(figsize=(15,8))
plt.plot([0, 1], [0, 1], linestyle='--')
# plotting ROC curve for various models
```

```
plt.plot(lr_fpr, lr_tpr, marker='.',c='r',label='LR Model without weights')
plt.plot(lrw_fpr, lrw_tpr, marker='.',c='y',label='LR Model with balanced_
↳weights')
plt.plot(lrwm_fpr, lrwm_tpr, marker='.',c='b',label='LR Model with manual_
↳weights')
plt.plot(dt_fpr, dt_tpr, marker='.',c='g',label='Decision tree Model')
plt.plot(rf_fpr, rf_tpr, marker='.',c='c',label='Random forest Model')
plt.plot(svc_fpr, svc_tpr, marker='.',c='m',label='SVC model')
plt.plot(knn_fpr, knn_tpr, marker='.',c='k',label='KNN model')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc='lower right')
plt.show()
```

ROC curves for various models



### 0.2.9 Calculating Sensitivity and Specificity for various models

```
[91]: def get_confusion_matrix_values(y_true, y_pred):
    cm = confusion_matrix(y_true, y_pred)
    return(cm[0][0], cm[0][1], cm[1][0], cm[1][1])

lr_TP, lr_FP, lr_FN, lr_TN = get_confusion_matrix_values(y, lr_model.predict(x))
lrw_TP, lrw_FP, lrw_FN, lrw_TN = get_confusion_matrix_values(y, lrw_model.
↳predict(x))
```

```

lrwm_TP, lrwm_FP, lrwm_FN, lrwm_TN = get_confusion_matrix_values(y, lrwm_model.
    ↪predict(x))
dt_TP, dt_FP, dt_FN, dt_TN = get_confusion_matrix_values(y, dt_model.predict(x))
rf_TP, rf_FP, rf_FN, rf_TN = get_confusion_matrix_values(y, rf_model.predict(x))
svc_TP, svc_FP, svc_FN, svc_TN = get_confusion_matrix_values(y, svc_model.
    ↪predict(x))
knn_TP, knn_FP, knn_FN, knn_TN = get_confusion_matrix_values(y, knn_model.
    ↪predict(x))

```

```

[92]: print('SENSITIVITY: \n')
print('Sensitivity of Logistic Regression model : ',(lr_TP/(lr_TP+lr_FN)))
print('Sensitivity of Logistic Regression model with balanced weight :␣
    ↪',(lrw_TP/(lrw_TP+lrw_FN)))
print('Sensitivity of Logistic regression model with manual weight : ',(lrwm_TP/
    ↪(lrwm_TP+lrwm_FN)))
print('Sensitivity of Decision Tree model : ',(dt_TP/(dt_TP+dt_FN)))
print('Sensitivity of Random forest model : ',(rf_TP/(rf_TP+rf_FN)))
print('Sensitivity of SVC model : ',(svc_TP/(svc_TP+svc_FN)))
print('Sensitivity of KNN model : ',(knn_TP/(knn_TP+knn_FN)))

```

SENSITIVITY:

```

Sensitivity of Logistic Regression model : 0.7876106194690266
Sensitivity of Logistic Regression model with balanced weight :
0.834070796460177
Sensitivity of Logistic regression model with manual weight :
0.8355263157894737
Sensitivity of Decision Tree model : 0.9545454545454546
Sensitivity of Random forest model : 0.9434697855750487
Sensitivity of SVC model : 0.7654941373534339
Sensitivity of KNN model : 0.8185255198487713

```

```

[93]: print('\nSPECIFICITY: \n')
print('Specificity of Logistic Regression model : ',(lr_TN/(lr_TN+lr_FP)))
print('Specificity of Logistic Regression model with balanced weight :␣
    ↪',(lrw_TN/(lrw_TN+lrw_FP)))
print('Specificity of Logistic regression model with manual weight : ',(lrwm_TN/
    ↪(lrwm_TN+lrwm_FP)))
print('Specificity of Decision Tree model : ',(dt_TN/(dt_TN+dt_FP)))
print('Specificity of Random forest model : ',(rf_TN/(rf_TN+rf_FP)))
print('Specificity of SVC model : ',(svc_TN/(svc_TN+svc_FP)))
print('Specificity of KNN model : ',(knn_TN/(knn_TN+knn_FP)))

```

SPECIFICITY:

```

Specificity of Logistic Regression model : 0.729064039408867

```

Specificity of Logistic Regression model with balanced weight :  
0.6107594936708861  
Specificity of Logistic regression model with manual weight :  
0.6185897435897436  
Specificity of Decision Tree model : 0.9351145038167938  
Specificity of Random forest model : 0.9372549019607843  
Specificity of SVC model : 0.7485380116959064  
Specificity of KNN model : 0.7196652719665272

#### **0.2.10 Various models are analysed and compared using different criterias and below are the findings**

##### **0.2.11 Based on accuracy of models on test data:**

Decision Tree model has highest accuracy of almost 74%. KNN model has an accuracy of 69% and so does Logistic Regression model.

##### **0.2.12 Based on ROC curve and AUC:**

Random Forest model shows best results followed by Decision Tree model, KNN model is far behind.

##### **0.2.13 Based on Sensitivity and Specificity:¶**

Decision Tree model has best sensitivity rate (95.45%) closely followed by Random forest (94.34%) while KNN model has a rate of 82.78%.

Specificity of Random Forest model is highest (93.72%) followed by Decision Tree model (93.51%). KNN model falls behind these two with a rate of 71.96%.

In this case we are building a model for disease detection and that's why we need to strictly minimize the number of False Negative hence we need maximum sensitivity in our model and based on this criteria we can choose Decision Tree model. The Random Forest model is performing Slightly better than any other model in all aspects except sensitivity, though it is very close to the top. Random Forest model can also be considered as it has almost equal sensitivity as Decision Tree model and far better specificity hence can reduce cost by reducing number of false positives

[ ]: